



THE UNIVERSITY  
*of* EDINBURGH

**SICSS Edinburgh**

# **Text Classification**

**Walid Magdy**

# Outlines

- Nature of text
  - Zipf's law
- Text Classification
  - Why you might need classification in CSS?
  - Feature Extraction
  - Feature Selection/Synthesis
  - Feature Weighting (e.g. TFIDF)
  - Classification process setup (train and test)
  - Evaluation

# Why Text Classification?

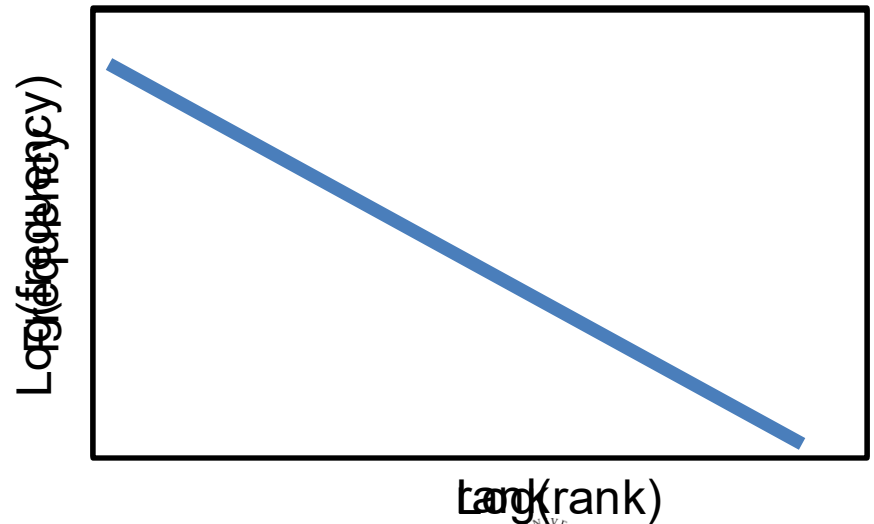
- Text → Most of social communication online .. *so far*
- CSS → Mostly analyse online data on large scale
- Data are not always labelled to be analysed on scale
- Some classifiers are available to use
  - E.g. Sentiment
- Sometimes you need to build a specific classifier for a certain task
  - E.g. Stance classifier for Pro-choice vs Pro-life
- Today: How to build a text classifier  
Tomorrow: How to build a general classifier

# Words' Nature

- Word → basic unit to represent text
- Certain characteristics are observed for the words we use!
- These characteristics are very consistent, that we can apply laws for them
- These laws apply for:
  - Different languages
  - Different domains of text

# Frequency of words

- Some words are very frequent  
e.g. “the”, “of”, “to”
- Many words are less frequent  
e.g. “schizophrenia”, “bazinga”
- ~50% terms appears once
- Frequency of words has hard exponential decay




# Zipf's Law:

- For a given collection of text, ranking unique terms according to their frequency, then:

$$r \times P_r \cong \text{const}$$

- $r$ , rank of term according to frequency
- $P_r$ , probability of appearance of term

- $P_r \cong \frac{\text{const}}{r} \rightarrow f(x) \cong \frac{1}{x}$



# Zipf's Law:

Wikipedia abstracts

→ 3.5M En abstracts

$$r \times P_r \cong \text{const} \rightarrow$$

$$r \times \text{freq}_r \cong \text{const}$$

Term	Rank	Frequency	r x freq
the	1	5,134,790	5,134,790
of	2	3,102,474	6,204,948
in	3	2,607,875	7,823,625
a	4	2,492,328	9,969,312
is	5	2,181,502	10,907,510
and	6	1,962,326	11,773,956
was	7	1,159,088	8,113,616
to	8	1,088,396	8,707,168
by	9	766,656	6,899,904
an	10	566,970	5,669,700
it	11	557,492	6,132,412
for	13	493,374	5,970,456
as	14	480,277	6,413,862
on	15	471,544	6,723,878
from	16	412,785	7,073,160

# Text Classification

- **Text classification** is the process of classifying documents into predefined categories based on their content.
  - Input: Text (document, article, sentence)
  - Task: Classify into predefined one/multiple categories
  - Categories:
    - Binary: relevant/irrelevant, spam .. etc.
    - Few: sports/politics/comedy/technology
    - Hierarchical: patents



# Classification is and is not

- **Classification** (a.k.a. “**categorization**”): a common technology in data science; studied within pattern recognition, statistics, and machine learning.
- Definition:  
the activity of **predicting** to which among a **predefined finite** set of groups (“classes”, or “categories”) a data item belongs to
- Formulated as the task of generating a hypothesis (or “classifier”, or “model”)

$$h : D \rightarrow C$$

where  $D = \{\mathbf{x}_1, \mathbf{x}_2, \dots\}$  is a domain of data items and  
 $C = \{c_1, \dots, c_n\}$  is a finite set of classes (the **classification scheme**)

# Classification is and is not

- Different from clustering, where the groups (“clusters”) and their number are not known in advance
- The membership of a data item into a class must not be determinable with certainty
  - e.g., predicting whether a natural number belongs to *Prime* or *Non-Prime* is not classification
- In text classification, data items are
  - **Textual**: e.g., news articles, emails, sentences, queries, etc.
  - **Partly textual**: e.g., Web pages

# Types of Classification

- **Binary:**

item to be classified into one of two classes

$$h : D \rightarrow C, \quad C = \{c_1, c_2\}$$

- e.g., Spam/not spam, offensive/not offensive, rel/irrel

- **Single-Label Multi-Class (SLMC)**

item to be classified into only one of  $n$  possible classes.

$$h : D \rightarrow C, \quad C = \{c_1 \dots c_n\}, \text{ where } n > 2$$

- e.g., Sports/politics/entertainment, positive/negative/neutral

- **Multi-Label Multi-Class (MLMC)**

item to be classified into none, one, two, or more classes

$$h : D \rightarrow 2^C, \quad C = \{c_1 \dots c_n\}, \text{ where } n > 1$$

- e.g., Assigning CS articles to classes in the ACM Classification System
- Usually be solved as  $n$  independent binary classification problems

# Dimension of Classification

- Text classification may be performed according to several dimensions (“axes”) orthogonal to each other.
- by **topic**; by far the most frequent case, its applications are global
- by **sentiment**; useful in market research, online reputation management, social science and political science.
- by **language** (e.g. “language identification”, “dialect identification”).
- by **genre**; e.g., AutomotiveNews vs. AutomotiveBlogs, useful in website classification and others;
- by **author** (a.k.a. “authorship attribution”), by native language (“native language identification”), or by gender; useful in forensics and cybersecurity
- by **usefulness**; e.g., product reviews

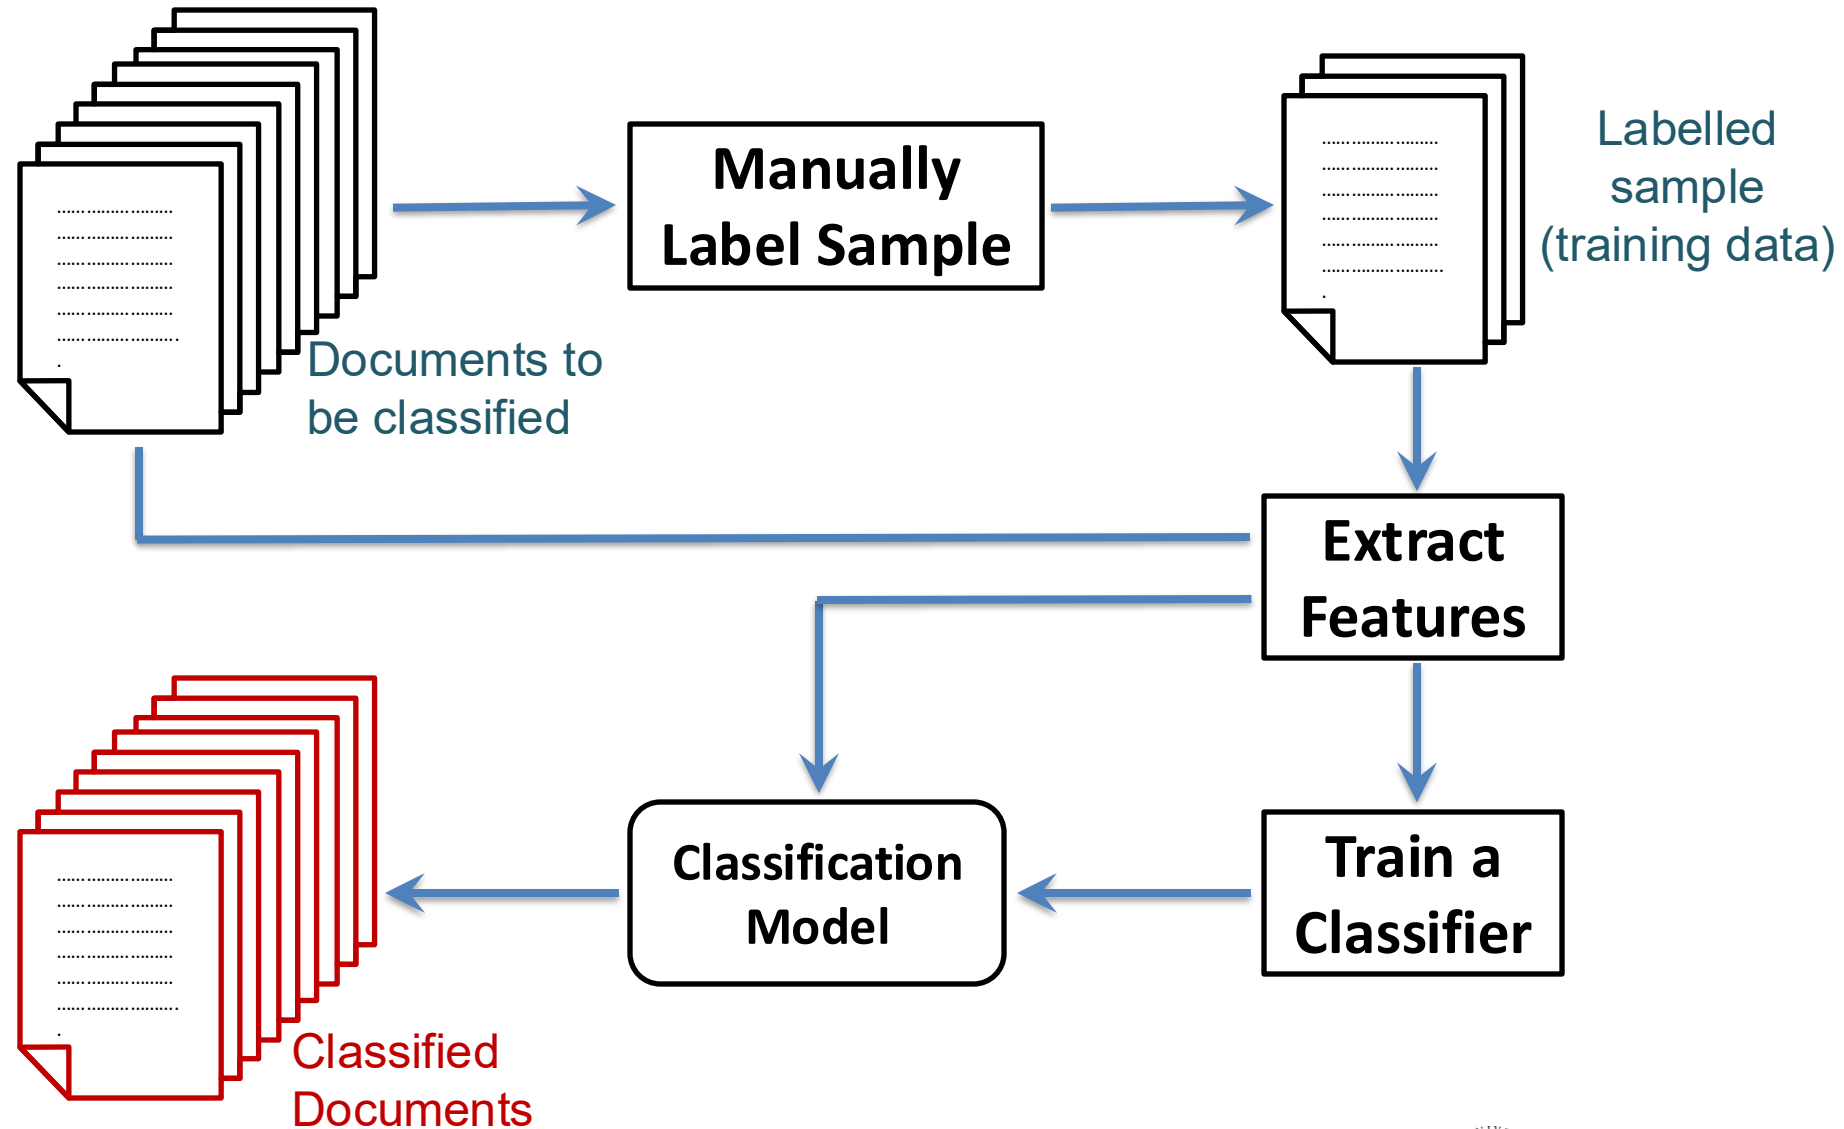
# Rule-based classification

- An old-fashioned way to build text classifiers was via knowledge engineering, i.e., manually building classification rules
  - E.g., (Viagra or Sildenafil or Cialis) → Spam
  - E.g. (#MAGA or America great again) → support Trump
- Common type: dictionary-based classification
- Disadvantages:
  - Expensive to setup and to maintain
  - Depends on few keywords → bad coverage (recall)
- Still accepted in many CSS studies
  - e.g. LIWC (Linguistic Inquiry and Word Count)

# Supervised-learning classification

- A generic (task-independent) learning algorithm is used to train a classifier from a set of manually classified examples
- The classifier learns, from these training examples, the characteristics a new text should have in order to be assigned to class  $c$
- Advantages:
  - Generating training examples cheaper than writing classification rules
  - Easy update to changing conditions (e.g., addition of new classes, deletion of existing classes, shifted meaning of existing classes, etc.)

# Supervised-learning classification



# Feature Extraction

- Objective: Text → Features set
- Steps: Extraction → Selection → Weighting
- Simplest form: BOW
  - Each term in a document is a feature
  - Feature space size = vocabulary in all docs
  - Standard preprocessing steps are usually applied
    - Tokenisation, stopping, stemming
- Other simple features forms:
  - Word n-grams (bigrams, trigrams, ....)
    - Much larger + more sparse
  - Sometimes char n-grams are used
    - Especially for degraded text (OCR or ASR outputs)



# Step 1: Feature Extraction

- What other text features could be used?
- Sentence structure (NLP):
  - POS (part-of-speech tags)
  - Syntactic tree structure
- Topic-based features (NLP):
  - LDA topics
  - NEs (named entities) in text
  - Links / Linked terms
- Non-textual features:
  - Average doc\sentence\word length
  - % of words start with upper-case letter
  - % of links/hashtags/emojis in text

# Step 1: Feature Extraction

- What preprocessing to apply?
  - Case-folding? **really** vs **Really** vs **REALLY**
  - Punctuation? “?”, “!”, “@”, “#”
  - Stopping? “**he**”, “**she**”, “**what**”, “**but**”
  - Stemming? “**replaced**” vs “**replacement**”
- Other Features:
  - Start with Cap, All Cap
  - Repeated characters “**congraaaaaats**” “**help!!!!!!!!!!**”
  - LIWC: Linguistic Inquiry and Word Count
- Which to choose?
  - Classification task/application

# Step 2: Feature Selection

- Number of distinctive features = feature space = length of feature vector.
- Vector can be of length  $O(10^6)$ , and might be sparse
  - High computational cost
  - Overfitting
- What are the most important features among those?
  - e.g. Reduce  $O(10^6)$  to  $O(10^4)$
- For each class, find the top representative  $k$  features for it → get the Union over all classes → reduced feature space

# Step 2: Feature Selection Functions

- Document frequency
  - % of docs in class  $c_i$  that contain the term  $t_k$
  - Very basic measure. Will select stop words as features

$$\#(t_k, c_i) = P(t_k | c_i)$$

- Mutual Information
  - How term  $t_k$  appear in class  $c_i$  compared to other classes
  - Highly used in feature selection in text classification

$$MI(t_k, c_i) = \sum_{c \in \{c_i, \bar{c}_i\}} \sum_{t \in \{t_k, \bar{t}_k\}} P(t, c) \cdot \log_2 \frac{P(t, c)}{P(t) \cdot P(c)}$$

- Pearson's Chi-squared ( $\chi^2$ )
  - used more in comparisons between classes

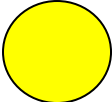
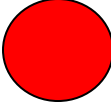
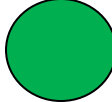
# Step 2: Feature Selection Functions


Function	Denoted by	Mathematical form
<i>Document frequency</i>	$\#(t_k, c_i)$	$P(t_k c_i)$
<i>DIA association factor</i>	$z(t_k, c_i)$	$P(c_i t_k)$
<i>Information gain</i>	$IG(t_k, c_i)$	$\sum_{c \in \{c_i, \bar{c}_i\}} \sum_{t \in \{t_k, \bar{t}_k\}} P(t, c) \cdot \log \frac{P(t, c)}{P(t) \cdot P(c)}$
<i>Mutual information</i>	$MI(t_k, c_i)$	$\log \frac{P(t_k, c_i)}{P(t_k) \cdot P(c_i)}$
<i>Chi-square</i>	$\chi^2(t_k, c_i)$	$\frac{ Tr  \cdot [P(t_k, c_i) \cdot P(\bar{t}_k, \bar{c}_i) - P(t_k, \bar{c}_i) \cdot P(\bar{t}_k, c_i)]^2}{P(t_k) \cdot P(\bar{t}_k) \cdot P(c_i) \cdot P(\bar{c}_i)}$
<i>NGL coefficient</i>	$NGL(t_k, c_i)$	$\frac{\sqrt{ Tr  \cdot [P(t_k, c_i) \cdot P(\bar{t}_k, \bar{c}_i) - P(t_k, \bar{c}_i) \cdot P(\bar{t}_k, c_i)]}}{\sqrt{P(t_k) \cdot P(\bar{t}_k) \cdot P(c_i) \cdot P(\bar{c}_i)}}$
<i>Relevancy score</i>	$RS(t_k, c_i)$	$\log \frac{P(t_k c_i) + d}{P(\bar{t}_k \bar{c}_i) + d}$
<i>Odds Ratio</i>	$OR(t_k, c_i)$	$\frac{P(t_k c_i) \cdot (1 - P(t_k \bar{c}_i))}{(1 - P(t_k c_i)) \cdot P(t_k \bar{c}_i)}$
<i>GSS coefficient</i>	$GSS(t_k, c_i)$	$P(t_k, c_i) \cdot P(\bar{t}_k, \bar{c}_i) - P(t_k, \bar{c}_i) \cdot P(\bar{t}_k, c_i)$

# Step 3: Feature Weighting

- Attributing a value to feature  $t_k$  in document  $d_i$   
This value may be
  - **binary** (representing presence/absence of  $t_k$  in  $d_i$ );
  - **numeric** (representing the importance of  $t_k$  for  $d_i$ );  
obtained via feature weighting functions in the following two classes:
    - **unsupervised**: e.g., tfidf or BM25,
    - **supervised**: e.g.,  $tf * MI$ ,  $tf * x^2$
- Current deep learning techniques, feature selecting and weighting is done automatically

# Should terms be weighted the same?

- Collection of 5 documents (balls = terms)
- Query   
- Which is the least relevant document?
- Which is the most relevant document?



# TFIDF

- **TFIDF:**  
Term Frequency, Inverse Document Frequency
- **$tf(t,d)$ :**  
number of times term  $t$  appeared in document  $d$ 
  - As  $tf(t,d) \uparrow\uparrow \rightarrow$  importance of  $t$  in  $d \uparrow\uparrow$
  - Document about IR, contains “retrieval” more than others
- **$df(t)$ :**  
number of documents term  $t$  appeared in
  - As  $df(d) \uparrow\uparrow \rightarrow$  importance if  $t$  in a collection  $\downarrow\downarrow$ 
    - “the” appears in many document  $\rightarrow$  not important
    - “FT” is not important word in financial times articles



# DF, CF, & IDF

- **DF  $\neq$  CF** (collection frequency)
  - $cf(t)$  = total number of occurrences of term  $t$  in a collection
  - $df(t) \leq N$  ( $N$ : number of documents in a collection)
  - $cf(t)$  can be  $\geq N$
- **DF** is more commonly used in IR than **CF**
  - **CF** is still used
- **$idf(t)$** : inverse of  **$df(t)$** 
  - As  $idf(t) \uparrow\uparrow \rightarrow$  rare term  $\rightarrow$  importance  $\uparrow\uparrow$
  - **$idf(t)$**   $\rightarrow$  measure of the informativeness of  $t$

# IDF: formula

$$idf(t) = \log_{10}\left(\frac{N}{df(t)}\right)$$

- ***idf(t)***: inverse of ***df(t)***
  - As  $idf(t) \uparrow\uparrow \rightarrow$  rare term  $\rightarrow$  importance  $\uparrow\uparrow$
  - ***idf(t)***  $\rightarrow$  measure of the informativeness of  $t$

- Suppose  $N = 1$  million  $\rightarrow$

term	<i>df(t)</i>	<i>idf(t)</i>
calpurnia	1	<b>6</b>
animal	100	<b>4</b>
sky	1,000	<b>3</b>
fly	10,000	<b>2</b>
under	100,000	<b>1</b>
the	1,000,000	<b>0</b>

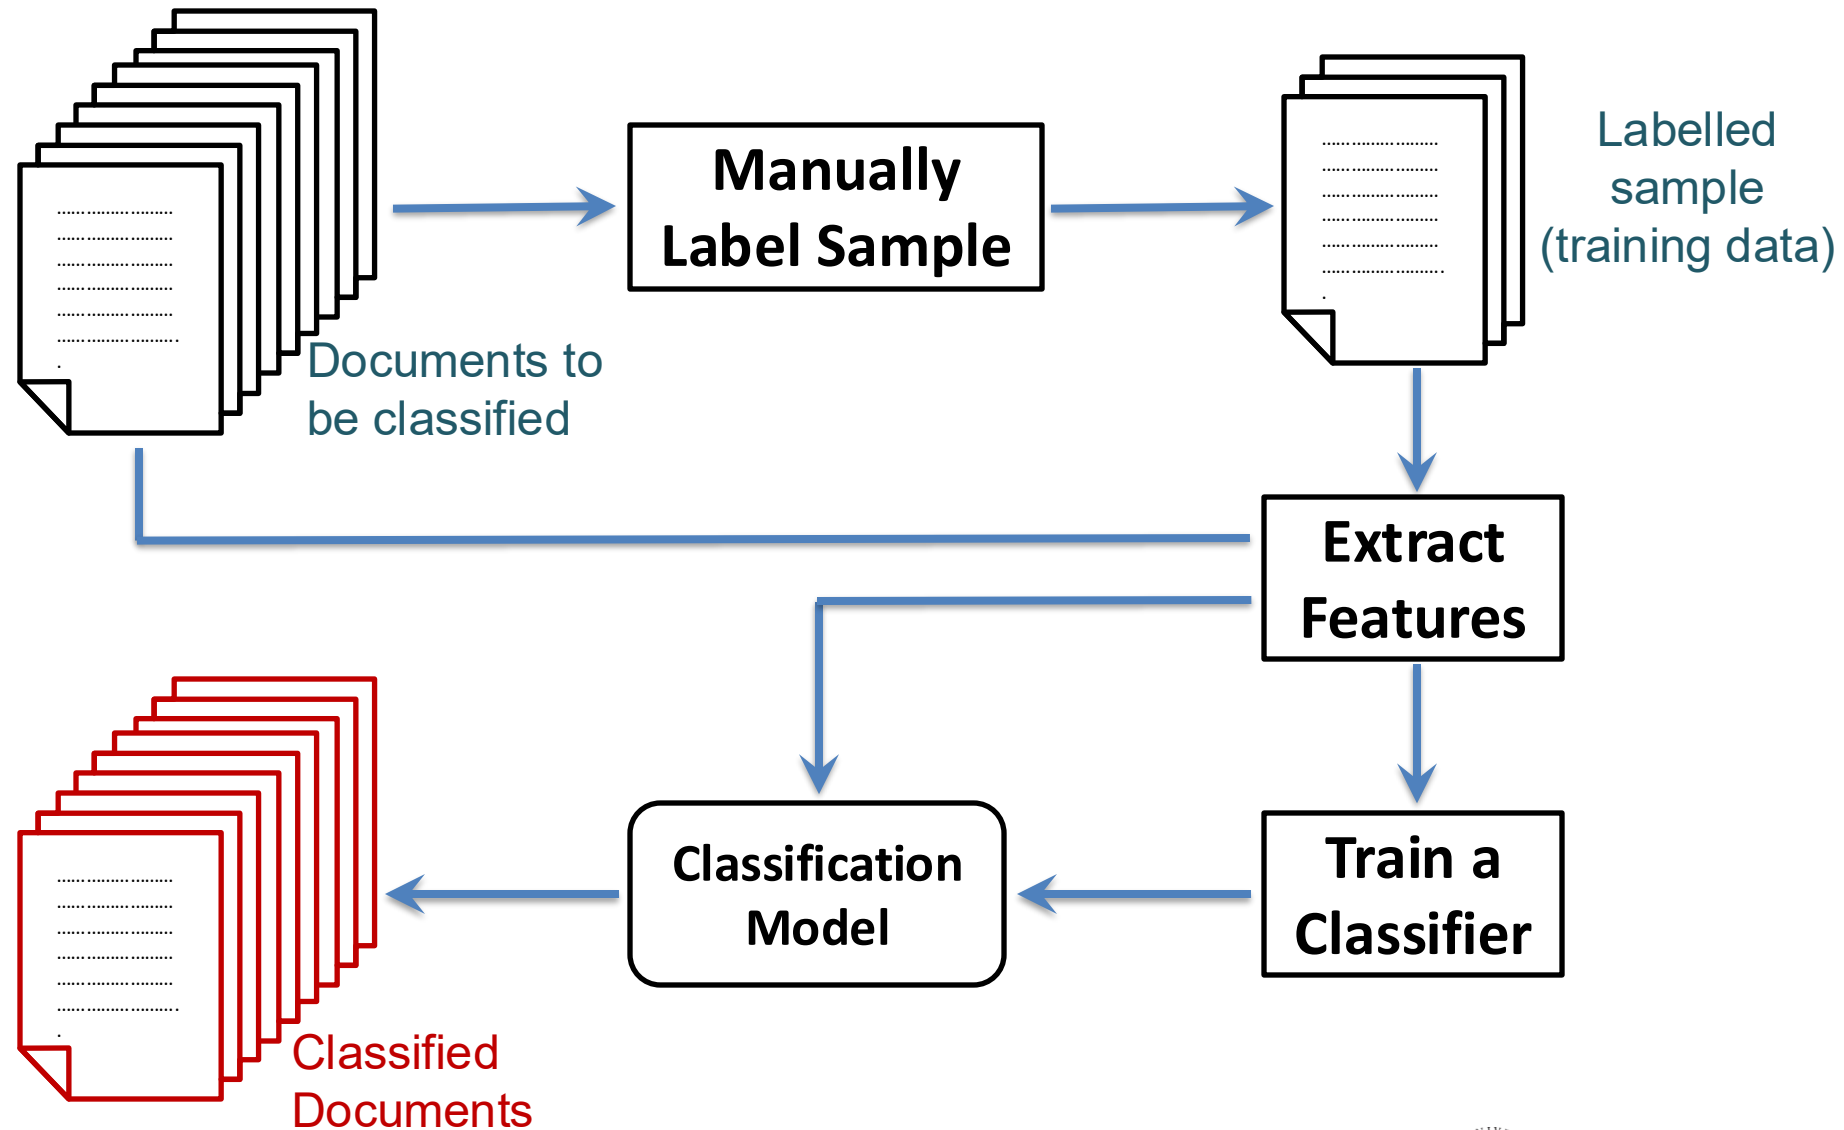
# TFIDF term weighting

- One the best known term weights schemes
  - Increases with the number of occurrences within a document
  - Increases with the rarity of the term in the collection
- Combines TF and IDF to find the weight of terms

$$w_{t.d} = \left(1 + \log_{10} tf(t, d)\right) \times \log_{10} \left(\frac{N}{df(t)}\right)$$

- With current ML techniques, new models learn term weight automatically

# Supervised-learning classification



# Training a Classifier

- For **binary** classification, essentially any supervised learning algorithm can be used for training a classifier; classical choices include
  - Support vector machines (SVMs)
  - Random forests
  - Naïve Bayesian methods
  - Lazy learning methods (e.g., k-NN)
  - Logistic Regression
  - ....
- The “**No-free-lunch principle**” (Wolpert, 1996) → *there is no learning algorithm that can outperform all others in all contexts*
- Implementations need to cater for
  - the very high dimensionality
  - the sparse nature of the representations involved

# Training a Classifier

- For **Multiclass classification**, some learning algorithms for binary classification are “SLMC-ready”; e.g.
  - Decision trees
  - Random forests
  - Naive Bayesian methods
  - Lazy learning methods (e.g., k-NN)
  - Neural networks
- For other learners (notably: SVMs) to be used for SLMC classification, combinations / cascades of the binary versions need to be used
  - e.g. multi-class classification SVM
  - Could be directly used for MLMC as well

# Parameter Optimisation of Classifier

- Most classifiers has some parameters to be optimized:  
(we will usually refer to the ones we set manually as “hyperparameters” to distinguish from the “learned” parameters/weights of the model)
  - The  $C$  parameter in soft-margin SVMs
  - The  $r, d$  parameters of non-linear kernels
  - Decision threshold for binary SVM
- Optimising the hyperparameters on test data is cheating!
- *Data Split*: Usually labelled data would be split into **three parts**
  - **Training**: used to train the classifier (typically **80%** of the data)
  - **Validation**: used to optimise hyperparameters. Apply the classifier on this data with different values of the hyperparameters and report the one that achieves the highest results (usually **10%** of the data)
  - **Test**: used to test the performance of the trained classifier with the optimal hyperparameters on these unseen data (usually **10%** of the data)

# Cross-Validation

- Sometimes the amount of labelled data in hand is limited (e.g. 200 samples). Having evaluation of a set of 20 samples only might be misleading
- Cross-validation is used to train the classifier with all data and test on all data without being cheating
- Idea:
  - Split the labelled data into  **$n$  folds**
  - Train classifier on  $n-1$  fold and test on the remaining one
  - Repeat  $n$  times
- **5-fold** cross validation
- Extreme case: LOOCV  
LOOCV: leave-one-out cross-validation



1
2
3
4
5



# State-of-the-art

- Transformer-based model
  - E.g.: BERT, RoBerta
  - No need for feature extraction/selection/weighting
  - No need for preprocessing
  - Excellent performance
  - Computationally more expensive (but light in inference)
- LLMs (generative models)
  - Zero-shot / few-shot
    - Sometimes not great without enough training
    - Extracting class from answer is sometimes not straightforward
  - Supervised Fine-Tuning (SFT)
    - Tune the model to do a specific task
    - Usually lead to good performance
  - Computationally very expensive (in training and inference)
- For classification tasks, when training data is available → transformer can be a more efficient and effective approach

# Evaluation

- Efficiency / Effectiveness
- Baselines
- Efficiency:
  - Training vs Inference (Classification)
  - Training data requirements
- Effectiveness:
  - Global effectiveness measures
  - Per class effectiveness measures

# Evaluation: Baselines




- There are standard methods for creating baselines

## Most popular/simplest baselines

- Random classification
  - Classes are assigned randomly
  - How much better is the classifier doing than random?
- Majority class baseline
  - Assign all elements to the class that appears the most
  - How much better you are doing than if you always picked the same thing output regardless of input?
- Simple algorithm, e.g. BOW
  - Usually used when you introduce new interesting features
- Recently: BERT baseline
- Now: zero-shot / few-shot baselines



# Evaluation: Binary Classification

- Accuracy:
  - How many of the samples are classified correctly?
- $A = 9/10 = 0.9$



# Evaluation: Binary Classification

- $A = 7/10 = 0.7$  System 1
- $A = 7/10 = 0.7$  System 2
- When classes are highly unbalanced
  - Precision/recall/F1 for the rare class
  - e.g. Spam classification (detection)



# Precision and Recall

- **Precision:**

What fraction of the classified as X are correct?

$$P = \frac{\textit{Classified correctly as X}}{\textit{All samples classified as X}}$$

- **Recall:**

What fraction of the class X has been classified correctly?


$$R = \frac{\textit{Classified correctly as X}}{\textit{Real number of the X samples}}$$



# F-measure

$$F1 = \frac{2 \cdot P \cdot R}{P + R}$$

- Harmonic mean of recall and precision
  - Emphasizes the importance of small values, whereas the arithmetic mean is affected more by outliers that are unusually large

# Evaluation: Binary Classification

	System 1	System 2
Precision	$1/3 = 0.33$	$0/1 = 0$
Recall	$1/2 = 0.5$	$0/2 = 0$
F1	0.4	0









# Evaluation: Multi-class

- Accuracy =  $(3+3+1)/10 = 0.7$
- Good measure when
  - Classes are nearly balanced
- Preferred:
  - Precision/recall/F1 for each class





			
P	0.75	1	0.333
R	0.75	0.75	0.5
F1	0.75	0.86	0.4

- **Macro-F1**  
 $= (0.75+0.86+0.4)/3$   
 $= \mathbf{0.67}$





# Evaluation: Multi-class

- Majority class baseline
- Accuracy = 0.8
- Macro-F1 = 0.296
- Macro-F1:
  - Should be used in binary classification when two classes are important
  - e.g.: males/females while distribution is 80/20%






# Error Analysis

- **Confusion Matrix**  
How classes get confused?

			
	3	0	1
	0	3	1
	1	0	1

- Useful:
  - Find classes that get confused with others
  - Develop better features to solve the problem



# Summary

- CSS requires classifying data for in-depth analysis
- Zipf's Law
- Text Classification tasks
- Feature extraction/selection/weighting
- Learning algorithms
- Cross-validation
- Baselines
- Evaluation measures
  - Accuracy/precision/recall/Macro-F1

# Resources

- *Fabrizio Sebastiani*  
**Machine Learning in Automated Text Categorization**  
*ACM Computing Surveys, 2002*  
Link: <https://arxiv.org/pdf/cs/0110053>
- *Yoav Goldberg*  
**A Primer on Neural Network Models for Natural Language Processing**  
Link: <https://arxiv.org/abs/1510.00726>
- *Vajjala, Sowmya, Shweta Shimangaud.*  
**Text Classification in the LLM Era--Where do we stand**  
Link: <https://arxiv.org/pdf/2502.11830>

# Practice

- **Zipf's distribution:**

<https://www.inf.ed.ac.uk/teaching/courses/tts/labs/lab1.html>

- **Text Classification**

<https://www.inf.ed.ac.uk/teaching/courses/tts/labs/lab7.html>

- Note: In-class practice tomorrow with Bjorn